

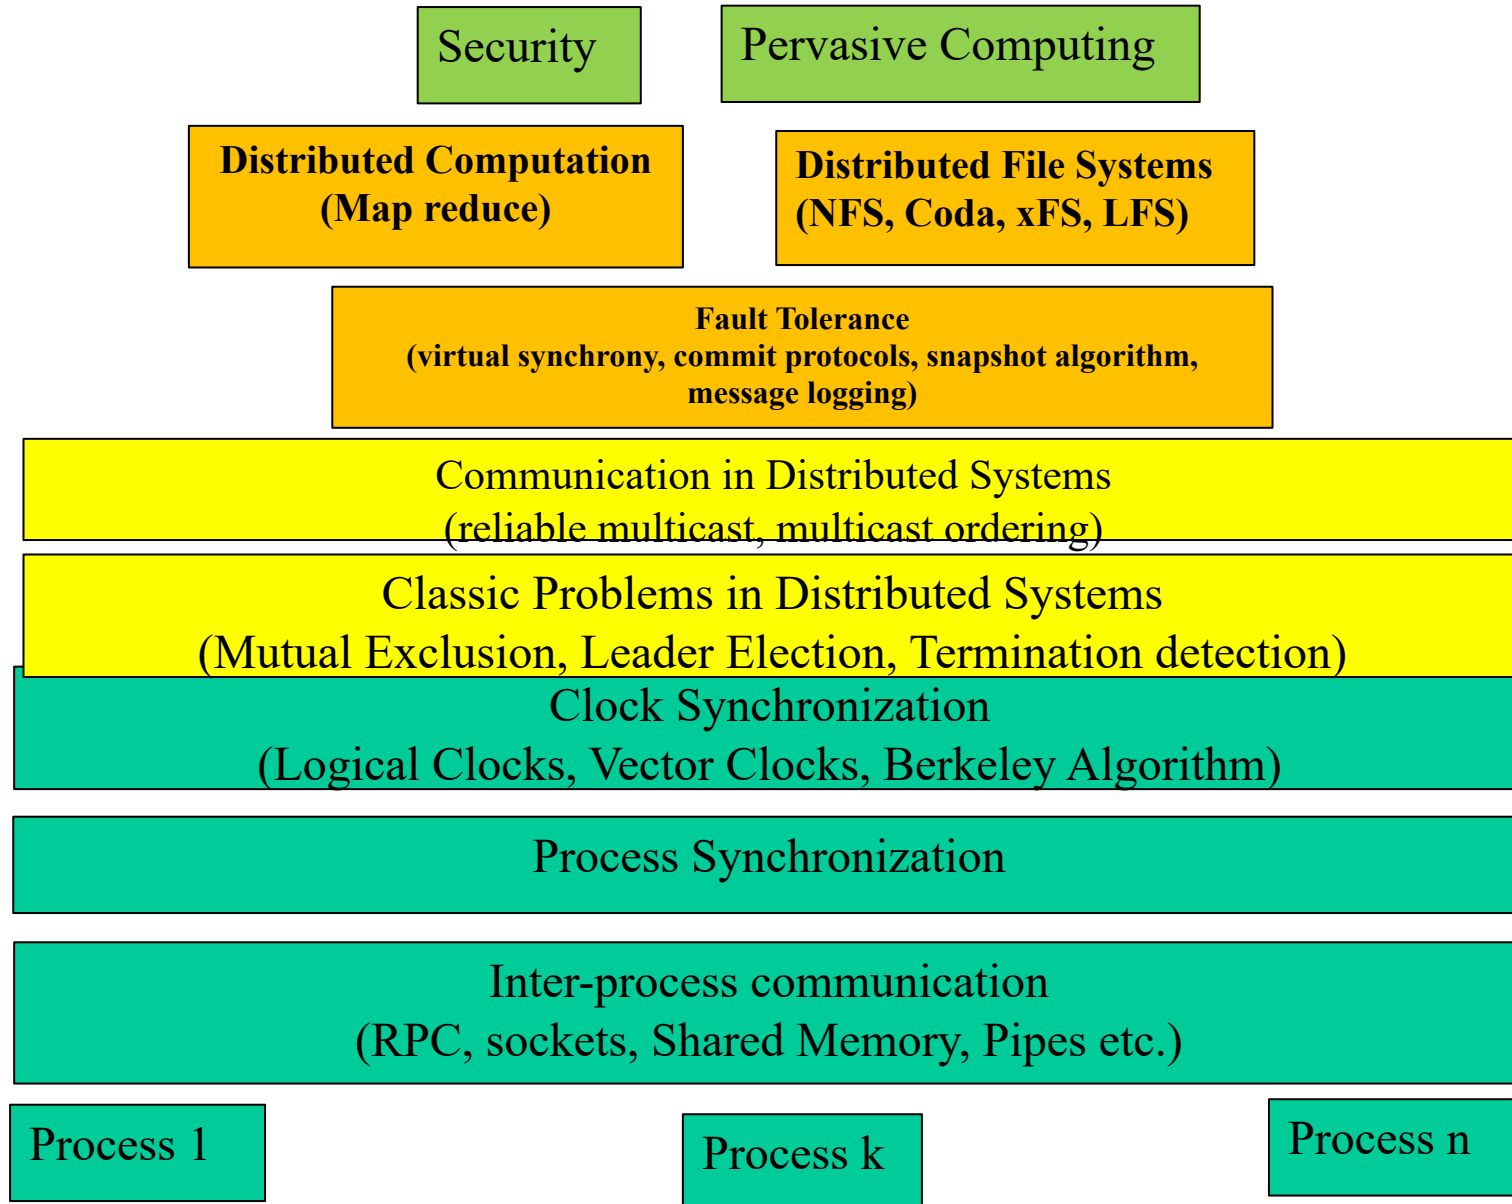
# CMSC621: Advanced Operating Systems

**Nilanjan Banerjee**

*Associate Professor, University of Maryland*  
Baltimore County  
nilanb@umbc.edu

<http://www.csee.umbc.edu/~nilanb/>


# Overview of the course



# Core Synchronization problems in Distributed Systems.


- Distributed Synchronization
  - Clocks, snapshots, leader election, distributed transactions

# Why is clock synchronization important?



When each machine has its own clock, an event that occurred after another event may nevertheless be assigned an earlier time.

## Another example: Global Positioning System



Computing a position in a two-dimensional space.

# Global Positioning System


Real world facts that complicate GPS

1. It takes a while before data on a satellite's position reaches the receiver.
2. The receiver's clock is generally not in synch with that of a satellite.

**What is the math behind calculating the GPS location?**

# Clock Synchronization Algorithms

The relation between clock time and UTC when clocks tick at different rates.





## Rate of Synchronization

If a manufacturer assures that the maximum clock drift is  $p$ . To assure that the relative offset of the clock is  $a$ , what the rate of synchronization or maximum interval of synchronization?


# Network Time Protocol

Getting the current time from a time server (Cristhian's algorithm)

Synchronize machines to a *time server* with a UTC receiver


Machine P requests time from server every  $\delta/2\rho$  seconds

- Receives time  $t$  from server, P sets clock to  $t + t_{reply}$  where  $t_{reply}$  is the time to send reply to P
- Use  $(t_{req} + t_{reply})/2$  as an estimate of  $t_{reply}$
- Improve accuracy by making a series of measurements




# The Berkeley Algorithm (1)

The time daemon asks all the other machines for their clock values.



## The Berkeley Algorithm (2)


The machines answer.



(b)

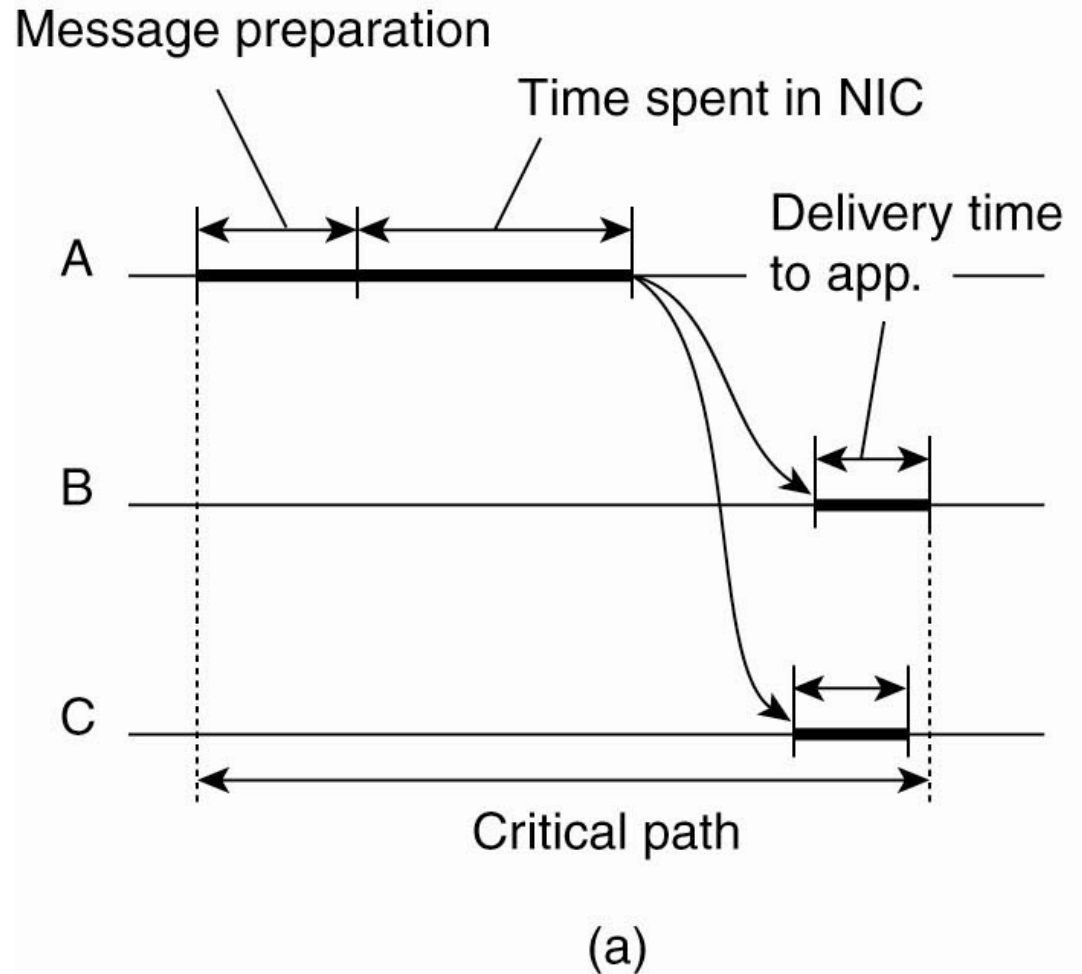
## The Berkeley Algorithm (3)

The time daemon tells everyone how to adjust their clock.



# Clock Synchronization in Wireless Networks

The usual critical path in determining network delays.



# Logical Clocks

- For many problems, internal consistency of clocks is important
  - Absolute time is less important
  - Use *logical* clocks
- Key idea:
  - Clock synchronization need not be absolute
  - If two machines do not interact, no need to synchronize them
  - More importantly, processes need to agree on the order in which events occur rather than the *time* at which they occurred

# Event Ordering

- *Problem:* define a total ordering of all events that occur in a system
- Events in a single processor machine are totally ordered
- In a distributed system:
  - No global clock, local clocks may be unsynchronized
  - Can not order events on different machines using local times
- Key idea [Lamport ]
  - Processes exchange messages
  - Message must be sent before received
  - Send/receive used to order events (and synchronize clocks)




# Lamport's Logical Clocks (1)

- The "happens-before" relation  $\rightarrow$  can be observed directly in two situations:
- If  $a$  and  $b$  are events in the same process, and  $a$  occurs before  $b$ , then  $a \rightarrow b$  is true.
- If  $a$  is the event of a message being sent by one process, and  $b$  is the event of the message being received by another process, then  $a \rightarrow b$
- What happens for processes that do not exchange messages?

## Lamport's Logical Clocks (2)


- (a) Three processes, each with its own clock. The clocks run at different rates.



(a)

Where is the happens-before relationship violated?


# Lamport's Logical Clocks



Lamport's algorithm corrects the clocks.

# Lamport's Logical Clocks

The positioning of Lamport's logical clocks in distributed systems.



# Lamport's Logical Clocks

Updating counter  $C_i$  for process  $P_i$


1. Before executing an event  $P_i$  executes  $C_i \leftarrow C_i + 1$ .
2. When process  $P_i$  sends a message  $m$  to  $P_j$ , it sets  $m$ 's timestamp  $ts(m)$  equal to  $C_i$  after having executed the previous step.
3. Upon the receipt of a message  $m$ , process  $P_j$  adjusts its own local counter as  $C_j \leftarrow \max\{C_j, ts(m)\} + 1$ , after which it then executes the first step and delivers the message to the application.

# Causality

- Lamport's logical clocks
  - If  $A \rightarrow B$  then  $C(A) < C(B)$ : **Proof?**
  - Is the reverse true?
    - Nothing can be said about events by comparing time-stamps!
    - If  $C(A) < C(B)$ , then what can you say?
- Need to maintain *causality*
  - If  $a \rightarrow b$  then  $a$  is causally related to  $b$
  - *Causal delivery*:
    - If  **$\text{send}(m) \rightarrow \text{send}(n) \Rightarrow \text{deliver}(m) \rightarrow \text{deliver}(n)$**
  - Capture causal relationships between groups of processes
  - Need a time-stamping mechanism such that:
    - If  $T(A) < T(B)$  then  $A$  should have causally preceded  $B$

## Vector Clocks (1)

Concurrent message transmission using logical clocks.



## Vector Clocks (2)

- Vector clocks are constructed by letting each process  $P_i$  maintain a vector  $VC_i$  with the following two properties:
  1.  $VC_i [ i ]$  is the number of events that have occurred so far at  $P_i$ . In other words,  $VC_i [ i ]$  is the local logical clock at process  $P_i$ .
  2. If  $VC_i [ j ] = k$  then  $P_i$  knows that  $k$  events have occurred at  $P_j$ . It is thus  $P_i$ 's knowledge of the local time at  $P_j$ .



## Vector Clocks (3)


- Steps carried out to accomplish property 2 of previous slide:
  1. Before executing an event  $P_i$  executes  $VC_i[i] \leftarrow VC_i[i] + 1.$
  2. When process  $P_i$  sends a message  $m$  to  $P_j$ , it sets  $m$ 's (vector) timestamp  $ts(m)$  equal to  $VC_i$  after having executed the previous step.
  3. Upon the receipt of a message  $m$ , process  $P_j$  adjusts its own vector by setting  $VC_j[k] \leftarrow \max\{VC_j[k], ts(m)[k]\}$  for each  $k$ , after which it executes the first step and delivers the message to the application.

## Software “Clock” 2: Vector Clocks

- Logical clock:
  - Event  $s$  happens before event  $t \Rightarrow$  the logical clock value of  $s$  is smaller than the logical clock value of  $t$ .
- Vector clock:
  - Event  $s$  happens before event  $t \Leftrightarrow$  the vector clock value of  $s$  is “smaller” than the vector clock value of  $t$ .
- Each event has a vector of  $n$  integers as its vector clock value
  - $v1 = v2$  if all  $n$  fields same
  - $v1 \leq v2$  if every field in  $v1$  is less than or equal to the corresponding field in  $v2$
  - $v1 < v2$  if  $v1 \leq v2$  and  $v1 \neq v2$

Relation “ $<$ ” here is  
not a total order

# Vector Clock Protocol



# Vector Clocks

**Lets do an in-class exercise**

# Vector Clock Properties

- Event  $s$  happens before  $t \Rightarrow$  vector clock value of  $s <$  vector clock value of  $t$ : There must be a chain from  $s$  to  $t$
- Event  $s$  happens before  $t \Leftarrow$  vector clock value of  $s <$  vector clock value of  $t$ 
  - If  $s$  and  $t$  on same process, done
  - If  $s$  is on  $p$  and  $t$  is on  $q$ , let  $V_s$  be  $s$ 's vector clock and  $V_t$  be  $t$ 's
  - $V_s < V_t \Rightarrow V_s[p] < V_t[p] \Rightarrow$  Must be a sequence of message from  $p$  to  $q$  after  $s$  and before  $t$

## Correctness of Vector Timestamps

**Theorem:** Vector timestamps implement vector clocks.

**Proof:** First, show  $a \rightarrow b$  implies

$$V(a) < V(b).$$

Case 1:  $a$  and  $b$  both occur at  $p_i$ ,  $a$  first. Since  $V_i$  increases at each step,

$$V(a) < V(b).$$

## Correctness of Vector Timestamps

Case 2:  $a$  occurs at  $p_i$  and causes  $m$  to be sent, while  $b$  occurs at  $p_j$  and includes the receipt of  $m$ .

- During  $b$ ,  $p_j$  updates its vector timestamp in such a way that  $V(a) \leq V(b)$ .
- $p_i$ 's estimate of number of steps taken by  $p_j$  is never an over-estimate. Since  $m$  is not received before it is sent,  $p_i$ 's estimate of the number of steps taken by  $p_j$  when  $a$  occurs is less than the number of steps taken by  $p_j$  when  $b$  occurs. So  $V(a)[j] < V(b)[j]$ .
- Thus  $V(a) < V(b)$ .

## Correctness of Vector Timestamps

Case 3: There exists  $c$  such that  $a \rightarrow c$  and  $c \rightarrow b$ .

By induction (from Cases 1 and 2) and transitivity of  $<$ ,  $V(a) < V(b)$ .

Next show  $V(a) < V(b)$  implies  $a \rightarrow b$ .

Equivalent to showing  $!(a \rightarrow b)$  implies  $!(V(a) < V(b))$



## Correctness of Vector Timestamps

- Suppose  $a$  occurs at  $p_i$ ,  $b$  occurs at  $p_j$ , and  $a$  does not happen before  $b$ .
- Let  $V(a)[i] = k$ .
- Since  $a$  does not happen before  $b$ , there is no chain of messages from  $p_i$  to  $p_j$  originating at  $p_i$  's  $k$ -th step or later and ending at  $p_j$  before  $b$ .
- Thus  $V(b)[i] < k$ .
- Thus  $\neg(V(a) < V(b))$ .

# Overview of the course

